



CSI3344 Distributed Systems

Workshop 04

- Q1. Assume a client calls an asynchronous RPC to a server, and subsequently waits until the server returns a result using another asynchronous RPC. Is this approach the same as letting the client execute a normal RPC?
- Q2. With persistent communication, a receiver generally has its own local buffer where messages can be stored when the receiver is not executing. To create such a buffer, we may need to specify its size. Give an argument why this is preferable, as well as one against specification of the size.
- Q3. Explain why transient synchronous communication has inherent scalability problems, and how these could be solved.

END OF THE WORKSHOP QUESTIONS